

Bank_pyspark

June 1, 2021

```
[ ]: # mount
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[ ]: !pip install pyspark
```

```
Collecting pyspark
  Downloading https://files.pythonhosted.org/packages/89/db/e18cfd78e408de
957821ec5ca56de1250645b05f8523d169803d8df35a64/pyspark-3.1.2.tar.gz (212.4MB)
    || 212.4MB 66kB/s
Collecting py4j==0.10.9
  Downloading https://files.pythonhosted.org/packages/9e/b6/6a4fb90cd235dc
8e265a6a2067f2a2c99f0d91787f06aca4bcf7c23f3f80/py4j-0.10.9-py2.py3-none-any.whl
(198kB)
    || 204kB 20.4MB/s
Building wheels for collected packages: pyspark
  Building wheel for pyspark (setup.py) ... done
  Created wheel for pyspark: filename=pyspark-3.1.2-py2.py3-none-any.whl
size=212880768
sha256=a9e7803006fd7770ea459f18acbb5247780943faf5ad09fd758870bcd3c5db89
  Stored in directory: /root/.cache/pip/wheels/40/1b/2c/30f43be2627857ab80062bef
1527c0128f7b4070b6b2d02139
Successfully built pyspark
Installing collected packages: py4j, pyspark
Successfully installed py4j-0.10.9 pyspark-3.1.2
```

```
[ ]: import pyspark
from pyspark.sql import SparkSession
spark = SparkSession.builder.appName("Bank-ml-pyspark").getOrCreate()

[ ]: df = spark.read.format('csv').option('header', True).load("/content/drive/
↳MyDrive/Bank_PySpark /bank.csv",inferSchema = True)

[ ]: df.show(5)
```

```

+---+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
|age|      job|marital|education|default|balance|housing|loan|contact|day|month
|duration|campaign|pdays|previous|poutcome|deposit|
+---+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
| 59|   admin.|married|secondary|   no|   2343|   yes|  no|unknown|  5|
may|   1042|     1|   -1|     0| unknown|   yes|
| 56|   admin.|married|secondary|   no|    45|   no|  no|unknown|  5|
may|   1467|     1|   -1|     0| unknown|   yes|
| 41|technician|married|secondary|   no|  1270|   yes|  no|unknown|  5|
may|   1389|     1|   -1|     0| unknown|   yes|
| 55|  services|married|secondary|   no|  2476|   yes|  no|unknown|  5|
may|    579|     1|   -1|     0| unknown|   yes|
| 54|   admin.|married| tertiary|   no|   184|   no|  no|unknown|  5|
may|    673|     2|   -1|     0| unknown|   yes|
+---+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
only showing top 5 rows

```

```
[ ]: df.printSchema()
```

```

root
|-- age: integer (nullable = true)
|-- job: string (nullable = true)
|-- marital: string (nullable = true)
|-- education: string (nullable = true)
|-- default: string (nullable = true)
|-- balance: integer (nullable = true)
|-- housing: string (nullable = true)
|-- loan: string (nullable = true)
|-- contact: string (nullable = true)
|-- day: integer (nullable = true)
|-- month: string (nullable = true)
|-- duration: integer (nullable = true)
|-- campaign: integer (nullable = true)
|-- pdays: integer (nullable = true)
|-- previous: integer (nullable = true)
|-- poutcome: string (nullable = true)
|-- deposit: string (nullable = true)

```

```
[ ]: df.groupby('deposit').count().show()
```

```

+-----+-----+
|deposit|count|

```

```

+-----+-----+
|      no| 5873|
|      yes| 5289|
+-----+-----+

```

```
[ ]: df.describe().show()
```

```

+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
|summary|          age|      job| marital|education|default|
balance|housing| loan| contact|          day|month|          duration|
campaign|          pdays|          previous|poutcome|deposit|
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
| count|          11162| 11162| 11162| 11162| 11162|
11162| 11162|11162| 11162|          11162|11162|          11162|
11162|          11162|          11162| 11162| 11162|
| mean|41.231947679627304| null| null| null|
null|1528.5385235620856| null| null| null|15.658036194230425|
null|371.99381831213043| 2.508421429851281|
51.33040673714388|0.8325568894463358| null| null|
| stddev|11.913369192215518| null| null| null| null|
3225.413325946149| null| null| null| 8.420739541006462|
null|347.12838571630687|2.7220771816614824|108.75828197197717|
2.292007218670508| null| null|
| min|          18| admin.|divorced| primary| no|
-6847| no| no|cellular|          1| apr|          2|
1|          -1|          0| failure| no|
| max|          95|unknown| single| unknown| yes|
81204| yes| yes| unknown|          31| sep|          3881|
63|          854|          58| unknown| yes|
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+

```

```
[ ]: df.dtypes
```

```
[ ]: [('age', 'int'),
      ('job', 'string'),
      ('marital', 'string'),
      ('education', 'string'),
      ('default', 'string'),
      ('balance', 'int'),
      ('housing', 'string'),
```

```
( 'loan', 'string'),
( 'contact', 'string'),
( 'day', 'int'),
( 'month', 'string'),
( 'duration', 'int'),
( 'campaign', 'int'),
( 'pdays', 'int'),
( 'previous', 'int'),
( 'poutcome', 'string'),
( 'deposit', 'string')]
```

```
[ ]: ## Extract all numeric columns
```

```
numeric_variables = []

for i in df.dtypes:
    if i[1] == 'int':
        numeric_variables.append(i[0])

numeric_variables
```

```
[ ]: ['age', 'balance', 'day', 'duration', 'campaign', 'pdays', 'previous']
```

```
[ ]: df.select(numeric_variables).describe().show()
```

```
+-----+-----+-----+-----+-----+
|summary|      age|    balance|      day|
duration|    campaign|      pdays|    previous|
+-----+-----+-----+-----+-----+
|  count|      11162|      11162|      11162|
11162|      11162|      11162|      11162|
|   mean|41.231947679627304|1528.5385235620856|15.658036194230425|371.9938183121
3043| 2.508421429851281| 51.33040673714388|0.8325568894463358|
|  stddev|11.913369192215518| 3225.413325946149|
8.420739541006462|347.12838571630687|2.7220771816614824|108.75828197197717|
2.292007218670508|
|   min|          18|        -6847|          1|
2|          1|          -1|          0|
|   max|          95|       81204|         31|
3881|          63|        854|          58|
+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
```

```
[ ]: numeric_data = df.select(numeric_variables)
numeric_data.show()
```

age	balance	day	duration	campaign	pdays	previous
59	2343	5	1042	1	-1	0
56	45	5	1467	1	-1	0
41	1270	5	1389	1	-1	0
55	2476	5	579	1	-1	0
54	184	5	673	2	-1	0
42	0	5	562	2	-1	0
56	830	6	1201	1	-1	0
60	545	6	1030	1	-1	0
37	1	6	608	1	-1	0
28	5090	6	1297	3	-1	0
38	100	7	786	1	-1	0
30	309	7	1574	2	-1	0
29	199	7	1689	4	-1	0
46	460	7	1102	2	-1	0
31	703	8	943	2	-1	0
35	3837	8	1084	1	-1	0
32	611	8	541	3	-1	0
49	-8	8	1119	1	-1	0
41	55	8	1120	2	-1	0
49	168	8	513	1	-1	0

only showing top 20 rows

```
[ ]: #Correlation Matrix
from pyspark.ml.feature import VectorAssembler
from pyspark.ml.stat import Correlation

assembler = VectorAssembler(
    inputCols = ['age', 'balance', 'day', 'duration', 'campaign', 'pdays', 'previous'],
    outputCol = "features"
)
assembled = assembler.transform(numeric_data)

pearson_corr = Correlation.corr(assembled, "features")

corr_list = pearson_corr.head()[0].toArray().tolist()
pearson_corr_df = spark.createDataFrame(corr_list)
pearson_corr_df.show(truncate=False)
```

_1	_2	_3	_4
----	----	----	----

```

|_5                |_6                |_7                |
+-----+-----+-----+-----+
|1.0                |0.11229988859873077
|-7.624209205460373E-4|1.892280737142355E-4
|-0.005277936156040414|0.002773834311769889|0.020168561218448653 |
|0.11229988859873077 |1.0                |0.010467439549070189
|0.022436131268962788 |-0.013893822542985367|0.01741114863267663
|0.03080524687156654 |
|-7.624209205460373E-4|0.010467439549070189 |1.0
|-0.018511399167089358|0.13700683429735389
|-0.07723161298141434|-0.05898068354621966 |
|1.892280737142355E-4 |0.022436131268962788 |-0.018511399167089358|1.0
|-0.04155745875962242 |-0.02739155324504362|-0.026716171271672622|
|-0.005277936156040414|-0.013893822542985367|0.13700683429735389
|-0.04155745875962242 |1.0
|-0.10272604750935362|-0.0496994979745621 |
|0.002773834311769889 |0.01741114863267663 |-0.07723161298141434
|-0.02739155324504362 |-0.10272604750935362 |1.0
|0.507271588372842    |
|0.020168561218448653 |0.03080524687156654 |-0.05898068354621966
|-0.026716171271672622|-0.0496994979745621 |0.507271588372842    |1.0
|
+-----+-----+-----+-----+
+-----+-----+-----+-----+

```

```
[ ]: df.show()
```

```

+---+-----+-----+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+-----+-----+
|age|      job| marital|education|default|balance|housing|loan|contact|day|mon
th|duration|campaign|pdays|previous|poutcome|deposit|
+---+-----+-----+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+-----+-----+
| 59|    admin.| married|secondary|    no|   2343|    yes|  no|unknown|  5|
may|   1042|      1|   -1|      0| unknown|    yes|
| 56|    admin.| married|secondary|    no|    45|    no|  no|unknown|  5|
may|   1467|      1|   -1|      0| unknown|    yes|
| 41| technician| married|secondary|    no|   1270|    yes|  no|unknown|  5|
may|   1389|      1|   -1|      0| unknown|    yes|
| 55|  services| married|secondary|    no|   2476|    yes|  no|unknown|  5|
may|    579|      1|   -1|      0| unknown|    yes|
| 54|    admin.| married| tertiary|    no|    184|    no|  no|unknown|  5|
may|    673|      2|   -1|      0| unknown|    yes|
| 42| management| single| tertiary|    no|     0|    yes| yes|unknown|  5|
may|    562|      2|   -1|      0| unknown|    yes|

```

56	management	married	tertiary	no	830	yes	yes	unknown	6
may	1201	1	-1	0	unknown	yes			
60	retired	divorced	secondary	no	545	yes	no	unknown	6
may	1030	1	-1	0	unknown	yes			
37	technician	married	secondary	no	1	yes	no	unknown	6
may	608	1	-1	0	unknown	yes			
28	services	single	secondary	no	5090	yes	no	unknown	6
may	1297	3	-1	0	unknown	yes			
38	admin.	single	secondary	no	100	yes	no	unknown	7
may	786	1	-1	0	unknown	yes			
30	blue-collar	married	secondary	no	309	yes	no	unknown	7
may	1574	2	-1	0	unknown	yes			
29	management	married	tertiary	no	199	yes	yes	unknown	7
may	1689	4	-1	0	unknown	yes			
46	blue-collar	single	tertiary	no	460	yes	no	unknown	7
may	1102	2	-1	0	unknown	yes			
31	technician	single	tertiary	no	703	yes	no	unknown	8
may	943	2	-1	0	unknown	yes			
35	management	divorced	tertiary	no	3837	yes	no	unknown	8
may	1084	1	-1	0	unknown	yes			
32	blue-collar	single	primary	no	611	yes	no	unknown	8
may	541	3	-1	0	unknown	yes			
49	services	married	secondary	no	-8	yes	no	unknown	8
may	1119	1	-1	0	unknown	yes			
41	admin.	married	secondary	no	55	yes	no	unknown	8
may	1120	2	-1	0	unknown	yes			
49	admin.	divorced	secondary	no	168	yes	yes	unknown	8
may	513	1	-1	0	unknown	yes			

```

+---+-----+-----+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+-----+-----+
only showing top 20 rows

```

```
[ ]: # Dropping day and month columns from original dataframe
```

```
df_final = df.drop("day", "month")
```

```
[ ]: df_final.show()
```

```

+---+-----+-----+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+-----+-----+
|age|      job| marital|education|default|balance|housing|loan|contact|duratio
n|campaign|pdays|previous|poutcome|deposit|
+---+-----+-----+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+-----+-----+
| 59|    admin.| married|secondary|    no|   2343|    yes|   no|unknown|
1042|      1|    -1|      0| unknown|   yes|
| 56|    admin.| married|secondary|    no|     45|    no|   no|unknown|
1467|      1|    -1|      0| unknown|   yes|

```

41	technician	married	secondary	no	1270	yes	no	unknown
1389	1	-1	0	unknown	yes			
55	services	married	secondary	no	2476	yes	no	unknown
579	1	-1	0	unknown	yes			
54	admin.	married	tertiary	no	184	no	no	unknown
673	2	-1	0	unknown	yes			
42	management	single	tertiary	no	0	yes	yes	unknown
562	2	-1	0	unknown	yes			
56	management	married	tertiary	no	830	yes	yes	unknown
1201	1	-1	0	unknown	yes			
60	retired	divorced	secondary	no	545	yes	no	unknown
1030	1	-1	0	unknown	yes			
37	technician	married	secondary	no	1	yes	no	unknown
608	1	-1	0	unknown	yes			
28	services	single	secondary	no	5090	yes	no	unknown
1297	3	-1	0	unknown	yes			
38	admin.	single	secondary	no	100	yes	no	unknown
786	1	-1	0	unknown	yes			
30	blue-collar	married	secondary	no	309	yes	no	unknown
1574	2	-1	0	unknown	yes			
29	management	married	tertiary	no	199	yes	yes	unknown
1689	4	-1	0	unknown	yes			
46	blue-collar	single	tertiary	no	460	yes	no	unknown
1102	2	-1	0	unknown	yes			
31	technician	single	tertiary	no	703	yes	no	unknown
943	2	-1	0	unknown	yes			
35	management	divorced	tertiary	no	3837	yes	no	unknown
1084	1	-1	0	unknown	yes			
32	blue-collar	single	primary	no	611	yes	no	unknown
541	3	-1	0	unknown	yes			
49	services	married	secondary	no	-8	yes	no	unknown
1119	1	-1	0	unknown	yes			
41	admin.	married	secondary	no	55	yes	no	unknown
1120	2	-1	0	unknown	yes			
49	admin.	divorced	secondary	no	168	yes	yes	unknown
513	1	-1	0	unknown	yes			

+---+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+
only showing top 20 rows

```
[ ]: print((df.count(), len(df.columns)))
      print((df_final.count(), len(df_final.columns)))
      print((numeric_data.count(), len(numeric_data.columns)))
      df_final.dtypes
```

(11162, 17)

(11162, 15)

(11162, 7)

```
[ ]: [('age', 'int'),
      ('job', 'string'),
      ('marital', 'string'),
      ('education', 'string'),
      ('default', 'string'),
      ('balance', 'int'),
      ('housing', 'string'),
      ('loan', 'string'),
      ('contact', 'string'),
      ('duration', 'int'),
      ('campaign', 'int'),
      ('pdays', 'int'),
      ('previous', 'int'),
      ('poutcome', 'string'),
      ('deposit', 'string')]
```

```
[ ]: from pyspark.ml.feature import OneHotEncoder, StringIndexer, VectorAssembler
categorical_variables = []

for i in df.dtypes:
    if i[1] == 'string':
        categorical_variables.append(i[0])
print(categorical_variables)
categorical_variables.remove('deposit')
categorical_variables.remove('month')
print(categorical_variables)
stages = []
print(stages)

#StringIndexer and OneHotEncoderEstimator for all categorical variables
for categorical_col in categorical_variables:
    stringIndexer = StringIndexer(inputCol= categorical_col, outputCol=
    →categorical_col + "Index")
    encoder = OneHotEncoder(inputCols=[stringIndexer.getOutputCol()],
    →outputCols=[categorical_col + "classVec"])
    stages += [stringIndexer, encoder]
print(stages)

#StringIndexer for target variable
string_indexer_label = StringIndexer(inputCol='deposit', outputCol= 'label')
stages +=[string_indexer_label]
print(stages)

#####Vector Assembler#####
```

```

###NOTE: When running the first time, please uncomment the next two commented
→code for removing the 2nd item in list (since we don't need it)
del(numeric_variables[2])
numeric_variables

assembler_inputs = [x + "classVec" for x in categorical_variables] +
    →numeric_variables
assembler_inputs
assembler = VectorAssembler(inputCols=assembler_inputs, outputCol="features")
stages +=[assembler]
print(stages)

```

```

['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact',
'month', 'poutcome', 'deposit']
['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact',
'poutcome']
[]

```

```

[StringIndexer_f9ec692fd80c, OneHotEncoder_6be121e2da2a,
StringIndexer_3aa7b56d57e1, OneHotEncoder_347aafc19a80,
StringIndexer_af6cf4689557, OneHotEncoder_ec6f0699a16f,
StringIndexer_651f1b4ea40f, OneHotEncoder_a8cd6d2c1a4f,
StringIndexer_621212989207, OneHotEncoder_57b800a9104c,
StringIndexer_7c75f1f74325, OneHotEncoder_e986f7099853,
StringIndexer_d95c52f9d115, OneHotEncoder_f3045f1ab21d,
StringIndexer_0d16e0841eb9, OneHotEncoder_58a373fc5b67]
[StringIndexer_f9ec692fd80c, OneHotEncoder_6be121e2da2a,
StringIndexer_3aa7b56d57e1, OneHotEncoder_347aafc19a80,
StringIndexer_af6cf4689557, OneHotEncoder_ec6f0699a16f,
StringIndexer_651f1b4ea40f, OneHotEncoder_a8cd6d2c1a4f,
StringIndexer_621212989207, OneHotEncoder_57b800a9104c,
StringIndexer_7c75f1f74325, OneHotEncoder_e986f7099853,
StringIndexer_d95c52f9d115, OneHotEncoder_f3045f1ab21d,
StringIndexer_0d16e0841eb9, OneHotEncoder_58a373fc5b67,
StringIndexer_414bf78dde83]
[StringIndexer_f9ec692fd80c, OneHotEncoder_6be121e2da2a,
StringIndexer_3aa7b56d57e1, OneHotEncoder_347aafc19a80,
StringIndexer_af6cf4689557, OneHotEncoder_ec6f0699a16f,
StringIndexer_651f1b4ea40f, OneHotEncoder_a8cd6d2c1a4f,
StringIndexer_621212989207, OneHotEncoder_57b800a9104c,
StringIndexer_7c75f1f74325, OneHotEncoder_e986f7099853,
StringIndexer_d95c52f9d115, OneHotEncoder_f3045f1ab21d,
StringIndexer_0d16e0841eb9, OneHotEncoder_58a373fc5b67,
StringIndexer_414bf78dde83, VectorAssembler_93d558470c57]

```

```

[ ]: #Building our Machine Learning Pipeline
cols = df_final.columns
from pyspark.ml import Pipeline

```

```

pipeline = Pipeline(stages = stages)
pipeline_model = pipeline.fit(df_final)
df_final = pipeline_model.transform(df_final)
all_columns = ['label', 'features'] + cols
df_final = df_final.select(all_columns)
df_final.dtypes

```

```

[:]: [ ('label', 'double'),
      ('features', 'vector'),
      ('age', 'int'),
      ('job', 'string'),
      ('marital', 'string'),
      ('education', 'string'),
      ('default', 'string'),
      ('balance', 'int'),
      ('housing', 'string'),
      ('loan', 'string'),
      ('contact', 'string'),
      ('duration', 'int'),
      ('campaign', 'int'),
      ('pdays', 'int'),
      ('previous', 'int'),
      ('poutcome', 'string'),
      ('deposit', 'string')]

```

```

[:]: df_pandas = df_final.toPandas()
      #df_pandas.head(3)
      df_pandas.iloc[3,1]

```

```

[:]: SparseVector(30, {4: 1.0, 11: 1.0, 13: 1.0, 16: 1.0, 18: 1.0, 20: 1.0, 21: 1.0,
24: 55.0, 25: 2476.0, 26: 579.0, 27: 1.0, 28: -1.0})

```

```

[:]: #df_pandas

```

```

[:]: categorical_variables
      for i in df_pandas[categorical_variables]:
          print(i, df_pandas[i].unique(), len(df_pandas[i].unique()))
      # We need to perform n - 1 for each category because of DropLast = True.
      # Currently, there are n categories for each feature in this loop.
      # Post n-1, we get a total of 24 categories in dataset and combining them with
      → 6 numeric features, gives us a total of 30 features, which can be
      # seen from the feature variable, which has been created via assembling all
      → data points per row to form a feature vector (per row).

```

```

job ['admin.' 'technician' 'services' 'management' 'retired' 'blue-collar'
     'unemployed' 'entrepreneur' 'housemaid' 'unknown' 'self-employed'
     'student'] 12
marital ['married' 'single' 'divorced'] 3
education ['secondary' 'tertiary' 'primary' 'unknown'] 4
default ['no' 'yes'] 2

```

```
housing ['yes' 'no'] 2
loan ['no' 'yes'] 2
contact ['unknown' 'cellular' 'telephone'] 3
poutcome ['unknown' 'other' 'failure' 'success'] 4
```

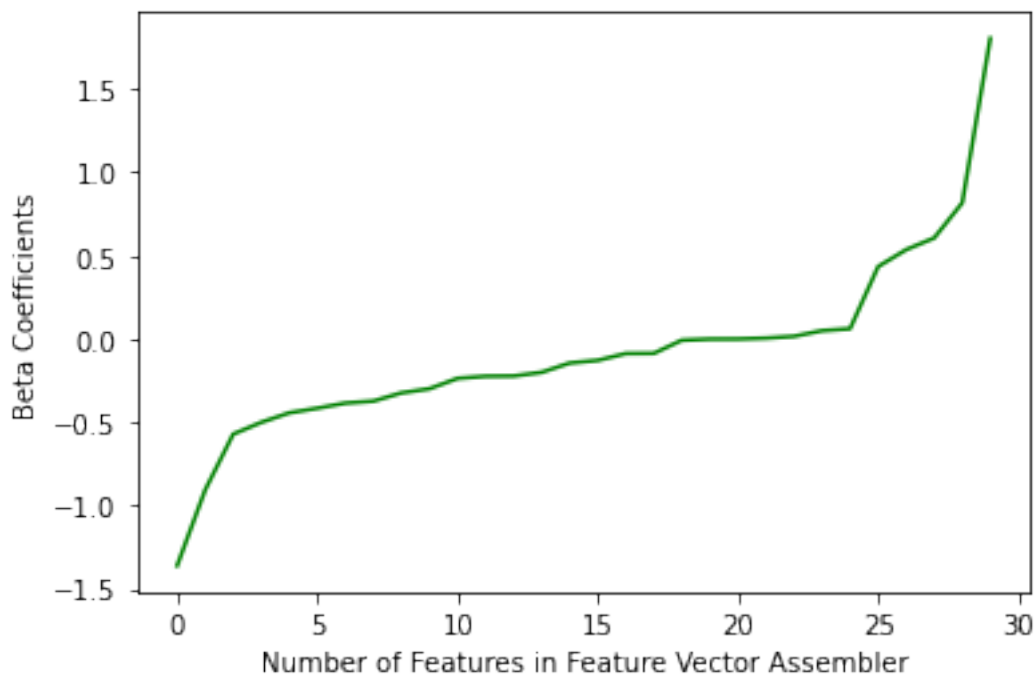
```
[ ]: #Train - Test Split
train, test = df_final.randomSplit([0.7, 0.3], seed = 111)
print("Training Dataset Count: " + str(train.count()))
print("Test Dataset Count: " + str(test.count()))
```

Training Dataset Count: 7843

Test Dataset Count: 3319

```
[ ]: #Logistic Regression
from pyspark.ml.classification import LogisticRegression
lr = LogisticRegression(featuresCol = 'features', labelCol = 'label',
    ↪maxIter=20)
lrModel = lr.fit(train)
```


```
[ ]: import matplotlib.pyplot as plt
import numpy as np
beta = np.sort(lrModel.coefficients)
plt.plot(beta, color = 'green')
plt.ylabel('Beta Coefficients')
plt.xlabel('Number of Features in Feature Vector Assembler')
plt.show()
```



```

[: trainingSummary = lrModel.summary
roc = trainingSummary.roc.toPandas()
plt.plot(roc['FPR'],roc['TPR'], color = "darkorange")
plt.ylabel('False Positive Rate')
plt.xlabel('True Positive Rate')
plt.title('ROC Curve')
plt.show()
print('Training set areaUnderROC: ' + str(trainingSummary.areaUnderROC))

```




Training set areaUnderROC: 0.8873176438519467

```

[: pr = trainingSummary.pr.toPandas()
plt.plot(pr['recall'],pr['precision'])
plt.ylabel('Precision')
plt.xlabel('Recall')
plt.show()

```



```
[ ]: predictions = lrModel.transform(test)
      predictions.select('age', 'job', 'label', 'rawPrediction', 'prediction', 'probability').show(10)
```

age	job	label	rawPrediction	prediction	probability
34	management	0.0	[0.23227426403389...	0.0	[0.55780889383587...
37	management	0.0	[1.07827232318531...	0.0	[0.74616689732780...
42	management	0.0	[1.42350053763192...	0.0	[0.80588660378317...
32	management	0.0	[1.13670611552191...	0.0	[0.75707436465925...
44	management	0.0	[0.94144140626849...	0.0	[0.71939072279361...
57	management	0.0	[1.08710009769855...	0.0	[0.74783525786982...
36	management	0.0	[1.13336142708195...	0.0	[0.75645870494901...
40	management	0.0	[1.53363719113977...	0.0	[0.82253785555589...
46	management	0.0	[2.10550235325418...	0.0	[0.89143682923626...
47	management	0.0	[0.60209164847146...	0.0	[0.64613469651098...

only showing top 10 rows

```
[ ]: from pyspark.ml.evaluation import BinaryClassificationEvaluator
      evaluator = BinaryClassificationEvaluator()
      print('Test Area Under ROC', evaluator.evaluate(predictions))
```

Test Area Under ROC 0.882180280221685

```
[ ]: # Decision Trees
from pyspark.ml.classification import DecisionTreeClassifier
dt = DecisionTreeClassifier(featuresCol = 'features', labelCol = 'label',
    ↳maxDepth = 3)
dtModel = dt.fit(train)
predictions = dtModel.transform(test)
predictions.select('age', 'job', 'label', 'rawPrediction', 'prediction',
    ↳'probability').show(10)
```

age	job	label	rawPrediction	prediction	probability
34	management	0.0	[845.0,1392.0]	1.0	[0.37773804202056...
37	management	0.0	[2489.0,487.0]	0.0	[0.83635752688172...
42	management	0.0	[2489.0,487.0]	0.0	[0.83635752688172...
32	management	0.0	[2489.0,487.0]	0.0	[0.83635752688172...
44	management	0.0	[2489.0,487.0]	0.0	[0.83635752688172...
57	management	0.0	[2489.0,487.0]	0.0	[0.83635752688172...
36	management	0.0	[416.0,38.0]	0.0	[0.91629955947136...
40	management	0.0	[416.0,38.0]	0.0	[0.91629955947136...
46	management	0.0	[2489.0,487.0]	0.0	[0.83635752688172...
47	management	0.0	[845.0,1392.0]	1.0	[0.37773804202056...

only showing top 10 rows

```
[ ]: evaluator = BinaryClassificationEvaluator()
print("Test Area Under ROC: " + str(evaluator.evaluate(predictions, {evaluator.
    ↳metricName: "areaUnderROC"})))
```

Test Area Under ROC: 0.7688262112991489

```
[ ]: from pyspark.ml.classification import RandomForestClassifier
rf = RandomForestClassifier(featuresCol = 'features', labelCol = 'label')
rfModel = rf.fit(train)
predictions = rfModel.transform(test)
predictions.select('age', 'job', 'label', 'rawPrediction', 'prediction',
    ↳'probability').show(10)
```

age	job	label	rawPrediction	prediction	probability
34	management	0.0	[9.60493129370130...	1.0	[0.48024656468506...
37	management	0.0	[14.7931681953967...	0.0	[0.73965840976983...

42 management	0.0	[14.2619205875643...	0.0	[0.71309602937821...
32 management	0.0	[15.1212713916369...	0.0	[0.75606356958184...
44 management	0.0	[14.7931681953967...	0.0	[0.73965840976983...
57 management	0.0	[14.5923319837989...	0.0	[0.72961659918994...
36 management	0.0	[14.0456675757228...	0.0	[0.70228337878614...
40 management	0.0	[14.0456675757228...	0.0	[0.70228337878614...
46 management	0.0	[16.5317096662812...	0.0	[0.82658548331406...
47 management	0.0	[8.74282839146477...	1.0	[0.43714141957323...

only showing top 10 rows

```
[ ]: evaluator = BinaryClassificationEvaluator()
print("Test Area Under ROC: " + str(evaluator.evaluate(predictions, {evaluator.
    ↳metricName: "areaUnderROC"})))
```

Test Area Under ROC: 0.8740465851804623

```
[ ]: #Gradient Boosting Tree Classifier
from pyspark.ml.classification import GBTClassifier
gbt = GBTClassifier(maxIter=10)
gbtModel = gbt.fit(train)
predictions = gbtModel.transform(test)
predictions.select('age', 'job', 'label', 'rawPrediction', 'prediction', '
    ↳probability').show(10)
```

age	job label	rawPrediction	prediction	probability
34 management	0.0	[-0.1379334480210...	1.0	[0.43146735146133...
37 management	0.0	[0.62197489330243...	0.0	[0.77625078341155...
42 management	0.0	[1.20862248237840...	0.0	[0.91813290106157...
32 management	0.0	[0.70697416510970...	0.0	[0.80438795213405...
44 management	0.0	[0.66751006881570...	0.0	[0.79166981202314...
57 management	0.0	[0.56219096811074...	0.0	[0.75480061571150...
36 management	0.0	[1.12218258958006...	0.0	[0.90416337754277...
40 management	0.0	[1.22491088090701...	0.0	[0.92054841562509...
46 management	0.0	[0.58568503698017...	0.0	[0.76339257374513...
47 management	0.0	[-0.2573502089919...	1.0	[0.37409228589155...

only showing top 10 rows

```
[ ]: evaluator = BinaryClassificationEvaluator()
print("Test Area Under ROC: " + str(evaluator.evaluate(predictions, {evaluator.
    ↳metricName: "areaUnderROC"})))
```

Test Area Under ROC: 0.8836635519598336


```
[ ]: # K-Fold Cross Validation with hyperparameter tuning (for GBT Classifier since
      ↳it provided highest accuracy prior to cross validation)
from pyspark.ml.tuning import ParamGridBuilder, CrossValidator
paramGrid = (ParamGridBuilder()
             .addGrid(gbt.maxDepth, [2, 4, 6])
             .addGrid(gbt.maxBins, [20, 60])
             .addGrid(gbt.maxIter, [10, 20])
             .build())
cv = CrossValidator(estimator=gbt, estimatorParamMaps=paramGrid,
                    ↳evaluator=evaluator, numFolds=5)

cvModel = cv.fit(train)
predictions = cvModel.transform(test)
evaluator.evaluate(predictions)
```

```
[ ]: 0.8910054649344297
```

```
[ ]: !wget -nc https://raw.githubusercontent.com/brpy/colab-pdf/master/colab_pdf.py
from colab_pdf import colab_pdf
colab_pdf('Bank_pyspark.ipynb')
```

File colab_pdf.py already there; not retrieving.

WARNING: apt does not have a stable CLI interface. Use with caution in scripts.

WARNING: apt does not have a stable CLI interface. Use with caution in scripts.

Extracting templates from packages: 100%

```
[ ]:
```